

GNU/Linux Performance tools: Profiling Rust



Couple of definitions

Benchmarking

<->

Performance measurements

<->

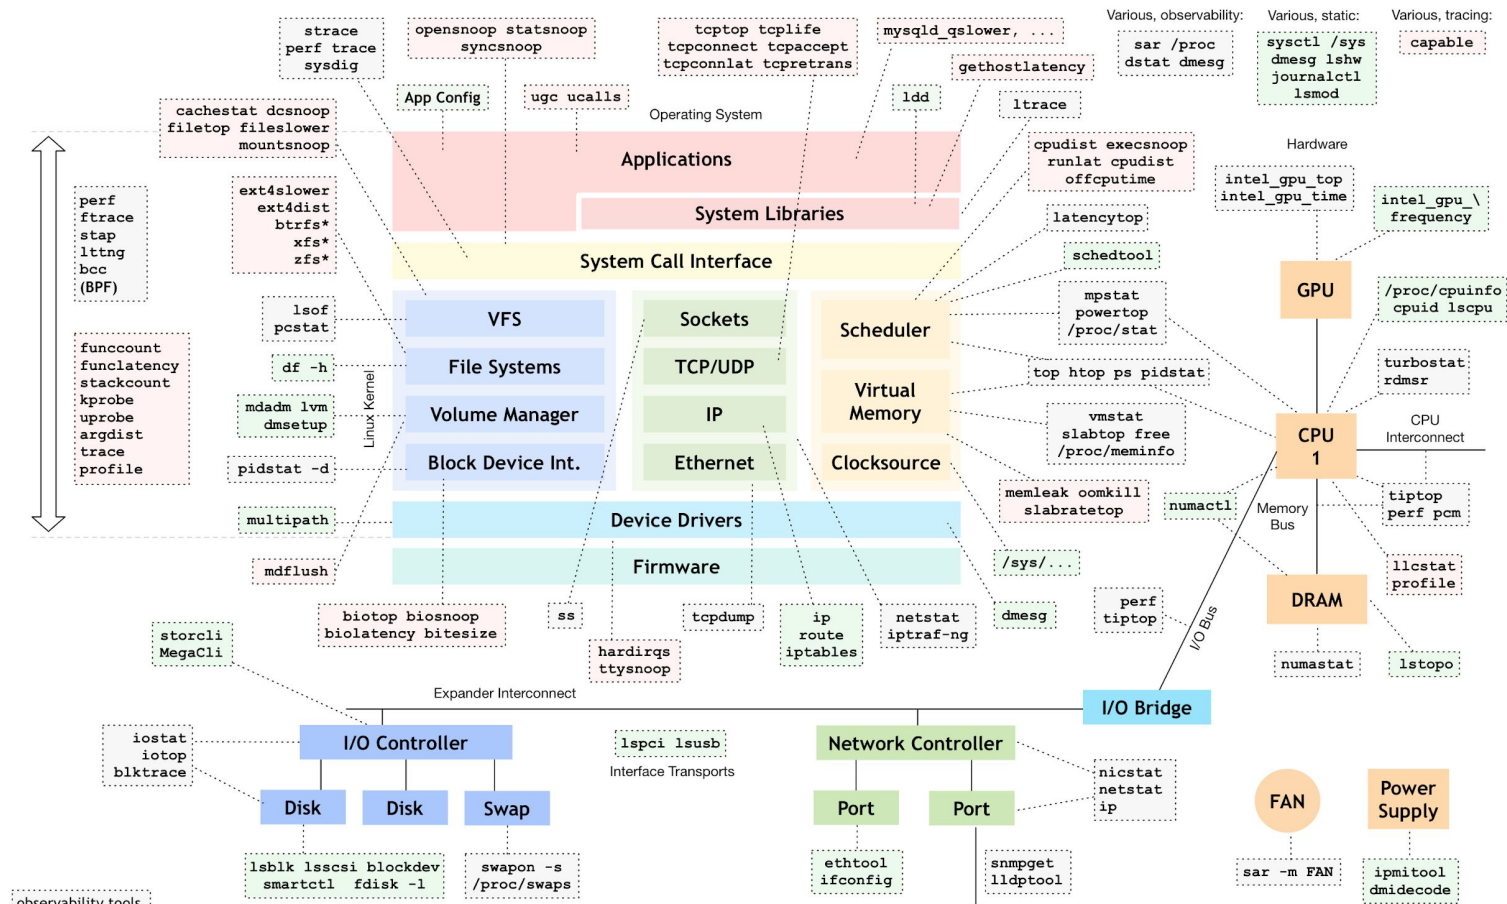
Tuning

<->

Configuration checks

Sampling <-> Counting

- Memory performance (memory allocation, memory leaks, memory stalls - eg fetches for cache)
- Lock contention (spin loops, blocking wait)
- CPU utilization
- Multithreading
- Disk IO
- Network IO
- ...



these can observe the state of the system at rest, without load

<https://github.com/brendangregg/perf-tools> <https://github.com/iovisor/bcc>

style inspired by reddit.com/u/redct

<http://www.brendangregg.com/linuxperf.html> 187

Poor man's profiling

- Run the binary (compile with debug symbols) to analyze under debugger
- Periodically stop the execution and get a backtrace
- Tools:
 - gdb
 - lldb
 - (VS Code)
 -

average: 0.03, 0.47, 0.33
ing, 0 stopped, 0 zombie

Real-time (interactive) list of linux processes.

- CPU summed across all CPUs, might lead to over 100%.

Usage:

- top

zxcVm1t0

A(w/a)

- htop

- 's' - trace processes' system calls (strace)
- 'l' - display open files for processes (lsOf)
- 'L' - trace library calls (ltrace) [only Debian]

- iotop

- per process/thread I/O monitor

SHR	S	%CPU	%MEM	TIME+	COMMAND
1376	S	0.0	0.0	0:00.58	init [4]
1300	S	0.0	0.0	0:01.01	- /sbin/udevd --daemon
2344	S	0.0	0.0		- /usr/sbin/cgmanager --daemon
1356	S	0.0	0.0		- /usr/sbin/syslogd
4	S	0.0	0.0		- /usr/sbin/klogd -c 3 -x
2184	S	0.0	0.0	0:00.13	- /usr/bin/dbus-daemon --system
1436	S	0.0	0.0		- /usr/sbin/inetd
4720	S	0.0	0.0	0:00.01	- /usr/sbin/ssh
4720	S	0.0	0.0	0:00.02	- /usr/sbin/acpid
4720	S	0.0	0.0	0:00.02	- /usr/sbin/console-kit-daemon

I	S	VIRT	RES	UID	COMMAND
0	S	4368	1468	0	init [4]
0	S	36100	3804	0	- /sbin/udevd --daemon
0	S	23768	2536	0	- /usr/sbin/cgmanager --daemon
0	S	6504	1480	0	- /usr/sbin/syslogd
0	S	4376	84	0	- /usr/sbin/klogd -c 3 -x
0	S	6688	224	0	- /sbin/dhccpd -t 10 eth0
0	S	20332	2800	81	- /usr/bin/dbus-daemon --system
0	S	6488	1540	0	- /usr/sbin/inetd
0	S	28676	428	0	- /usr/sbin/ssh
0	S	4392	120	0	- /usr/sbin/acpid
0	S	410392	7340	0	- /usr/sbin/console-kit-daemon

I	S	VIRT	RES	UID	COMMAND
6	S	1376	18	0	init [4]
2	S	2300		0	- /sbin/udevd --daemon
0	S	2344		0	- /usr/sbin/cgmanager --daemon
4	S	1356		0	- /usr/sbin/syslogd
4	S			0	- /usr/sbin/klogd -c 3 -x
2	S	4		0	- /sbin/dhccpd -t 10 eth0
6	S	2184		0	- /usr/bin/dbus-daemon --system
4	S	1436		0	- /usr/sbin/inetd
4	S			0	- /usr/sbin/ssh
4	S			0	- /usr/sbin/acpid
0	S	4720		0	- /usr/sbin/console-kit-daemon

TIME+	%CPU	%MEM	S	COMMAND
0:00.58	0.0	0.0	S	init [4]
0:01.01	0.0	0.0	S	- /sbin/udevd --daemon
	0.0	0.0	S	- /usr/sbin/cgmanager --daemon
	0.0	0.0	S	- /usr/sbin/syslogd
	0.0	0.0	S	- /usr/sbin/klogd -c 3 -x
	0.0	0.0	S	- /sbin/dhccpd -t 10 eth0
0:00.13	0.0	0.0	S	- /usr/bin/dbus-daemon --system
	0.0	0.0	S	- /usr/sbin/inetd
	0.0	0.0	S	- /usr/sbin/ssh
	0.0	0.0	S	- /usr/sbin/acpid
0:00.02	0.0	0.0	S	- /usr/sbin/console-kit-daemon
0:00.02	0.0	0.0	S	- /usr/lib/polkit-1/polkitd --no-debug
0:00.05	0.1	0.0	S	- /usr/lib/polkit-1/polkitd --no-debug

```
PIDSTAT(1)                                Linux User's Manual                                PIDSTAT(1)

NAME
    pidstat - Report statistics for Linux tasks

SYNOPSIS
    pidstat [-d] [-H] [-h] [-I] [-l] [-R] [-r] [-s] [-t] [-U username]
    [-u] [-V] [-v] [-w] [-C comm] [-G process_name] [--human] [-p pid]
    [SELF | ALL] [-T TASK | CHILD | ALL] [interval [count]] [-e program args]

DESCRIPTION
    The pidstat command is used for monitoring individual tasks currently being managed by the Linux kernel. It writes to standard output activities for every task selected with option -p or for every task managed by the Linux kernel if option -p ALL has been used. Selecting any tasks is equivalent to specifying -p ALL but only active tasks (tasks with non-zero statistics values) will appear in the report.

    The pidstat command can also be used for monitoring the child processes of selected tasks. Read about option -T below.

    The interval parameter specifies the amount of time in seconds between each report. A value of 0 (or no parameters at all) indicates that tasks statistics are to be reported continuously since the time since system startup (boot). The count parameter can be specified in conjunction with the interval parameter if this one is not set to zero. The value of count determines the number of reports generated at interval seconds apart. If the interval parameter is specified without the count parameter, the pidstat command generates reports continuously.

    You can select information about specific task activities using flags. Not specifying any flags selects only CPU activity.

OPTIONS
    -C comm
        Display only tasks whose command name includes the string comm. This string can be a regular expression.

    -d
        Report I/O statistics (kernels 2.6.20 and later only). The following values may be displayed:

        UID
            The real user identification number of the task being monitored.

        USER
            The name of the real user owning the task being monitored.

        PID
            The identification number of the task being monitored.

        kB_rd/s
            Number of kilobytes the task has caused to be read from disk per second.

        kB_wr/s
            Number of kilobytes the task has caused to be written to disk per second.
```

Used for monitoring of individual tasks currently being managed by the kernel.

Usage:

- Memory utilization, page faults, stack utilization, CPU utilization, task switching,...
- d Report I/O statistics
- r Report page faults and memory utilization
- s Report stack utilization
- u Report CPU utilization
- w Report task switching activity

Sysstat family


iotat – Report CPU statics and input/output statistics for devices, partitions and network filesystems (NFS), eg monitoring IO on your dedicated disk.

mpstat – Report processors related statistics, use eg for monitoring of pinned processes.

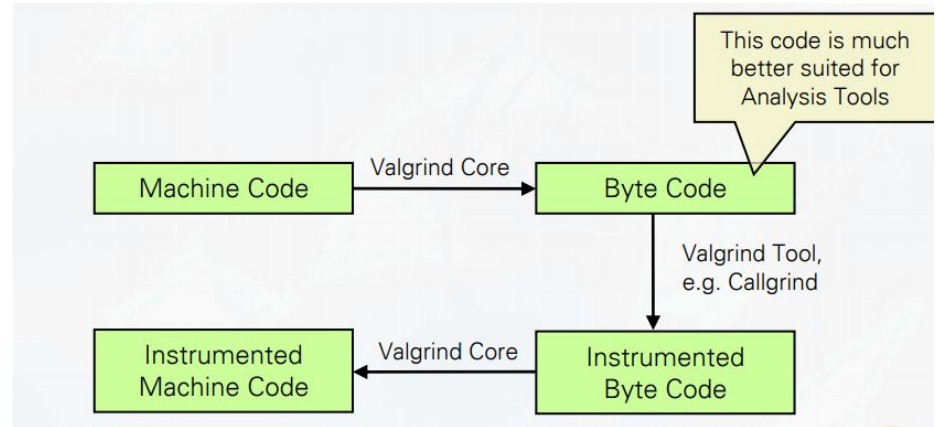
sar – Collect, report, or save system activity information.

...

valgrind tools



- Suite of simulation-based debugging and profiling tools
- program is run on a synthetic CPU provided by the Valgrind core
- Tools are adding own instrumentation code to binary to be handled by Valgrind's core
(Valgrind Tool = Valgrind Core + Tool Plugin)
- Large overhead (~10-100x slower)



valgrind tools



memcheck – memory error detection

cachegrind – cache and branch prediction profiler

callgrind – a callgraph generating cache and branch prediction profiler

massif – heap profiler

helgrind, DRD – thread error detectors
(experimental: stack/global array overrun detector, heap profiler – DHAT, that examines how heap blocks are used,..)

```

==3771== HEAP SUMMARY:
==3771== in use at exit: 0 bytes in 0 blocks
==3771== total heap size: 1,024 bytes in 4 blocks
==3771==
==3771== 32 bytes in 1 blocks are still reachable in loss record 1 of 4
==3771== at 0x128A12: je_malloc (in ./Programs/Rust/learn)
==3771== by 0x11137f: thread_local::imp::register dtor::hbd14515a39f84cf203b (in ./Programs/Rust/learn)
==3771== by 0x1129E3: thread_local::Key$LT$T$GT$::with::h7241079596358413660 (in ./Programs/Rust/learn)
==3771== by 0x112C42: io::stdio::with task stdout::hafb602ec3975fdc5syg (in ./Programs/Rust/learn)
==3771== by 0x115557: io::stdio::println args::h7f0794513d4fd2fcxDbg (in ./Programs/Rust/learn)
==3771== by 0x111087: main::hbd01322dca53aee9aaa (in ./Programs/Rust/learn)
==3771== by 0x11A22A: rt::start::closure.32122 (in ./Programs/Rust/learn)
==3771== by 0x124FEB: rust try inner (in ./Programs/Rust/learn)
==3771== by 0x124FD5: rust try (in ./Programs/Rust/learn)
==3771== by 0x1238B2: unwind::try::hf2f7fc7ecc46c43Tyc (in ./Programs/Rust/learn)
==3771== by 0x123A88: task::Task::run::h911f3b3bbb0c433efKb (in ./Programs/Rust/learn)
==3771== by 0x11A058: rt::start::hfd264fa826df360859x (in ./Programs/Rust/learn)
==3771==
==3771== 64 bytes in 1 blocks are still reachable in loss record 2 of 4
==3771== at 0x128A12: je_malloc (in ./Programs/Rust/learn)
==3771== by 0x111431: thread_local::imp::register dtor::hbd14515a39f84cf203b (in ./Programs/Rust/learn)
==3771== by 0x1129E3: thread_local::Key$LT$T$GT$::with::h7241079596358413660 (in ./Programs/Rust/learn)
==3771== by 0x112C42: io::stdio::with task stdout::hafb602ec3975fdc5syg (in ./Programs/Rust/learn)
==3771== by 0x115557: io::stdio::println args::h7f0794513d4fd2fcxDbg (in ./Programs/Rust/learn)
==3771== by 0x111087: main::hbd01322dca53aee9aaa (in ./Programs/Rust/learn)
==3771== by 0x11A22A: rt::start::closure.32122 (in ./Programs/Rust/learn)
==3771== by 0x124FEB: rust try inner (in ./Programs/Rust/learn)
==3771== by 0x124FD5: rust try (in ./Programs/Rust/learn)
==3771== by 0x1238B2: unwind::try::hf2f7fc7ecc46c43Tyc (in ./Programs/Rust/learn)
==3771== by 0x123A88: task::Task::run::h911f3b3bbb0c433efKb (in ./Programs/Rust/learn)
==3771== by 0x11A058: rt::start::hfd264fa826df360859x (in ./Programs/Rust/learn)
==3771==
==3771== 64 bytes in 1 blocks are still reachable in loss record 3 of 4
==3771== at 0x128A12: je_malloc (in ./Programs/Rust/learn)
==3771== by 0x112CD0: io::stdio::with task stdout::hafb602ec3975fdc5syg (in ./Programs/Rust/learn)
==3771== by 0x115557: io::stdio::println args::h7f0794513d4fd2fcxDbg (in ./Programs/Rust/learn)
==3771== by 0x111087: main::hbd01322dca53aee9aaa (in ./Programs/Rust/learn)
==3771== by 0x11A22A: rt::start::closure.32122 (in ./Programs/Rust/learn)
==3771== by 0x124FEB: rust try inner (in ./Programs/Rust/learn)
==3771== by 0x124FD5: rust try (in ./Programs/Rust/learn)
==3771== by 0x1238B2: unwind::try::hf2f7fc7ecc46c43Tyc (in ./Programs/Rust/learn)
==3771== by 0x123A88: task::Task::run::h911f3b3bbb0c433efKb (in ./Programs/Rust/learn)
==3771== by 0x11A058: rt::start::hfd264fa826df360859x (in ./Programs/Rust/learn)
==3771== by 0x119E35: rt::lang_start::hde2a214462357c7eb9x (in ./Programs/Rust/learn)
==3771== by 0x11117E: main (in ./Programs/Rust/learn)
==3771==
==3771== 1,024 bytes in 1 blocks are still reachable in loss record 4 of 4
==3771== at 0x128A12: je_malloc (in ./Programs/Rust/learn)
==3771== by 0x112222: io::buffered::Writer$LT$T$WT$::with capacity::h8206835914802320726 (in ./Programs/Rust/learn)
==3771== by 0x112063: io::stdio::with task stdout::hafb602ec3975fdc5syg (in ./Programs/Rust/learn)
==3771== by 0x115557: io::stdio::println args::h7f0794513d4fd2fcxDbg (in ./Programs/Rust/learn)
==3771== by 0x111087: main::hbd01322dca53aee9aaa (in ./Programs/Rust/learn)
==3771== by 0x11A22A: rt::start::closure.32122 (in ./Programs/Rust/learn)
==3771== by 0x124FEB: rust try inner (in ./Programs/Rust/learn)
==3771== by 0x124FD5: rust try (in ./Programs/Rust/learn)
==3771== by 0x1238B2: unwind::try::hf2f7fc7ecc46c43Tyc (in ./Programs/Rust/learn)
==3771== by 0x123A88: task::Task::run::h911f3b3bbb0c433efKb (in ./Programs/Rust/learn)
==3771== by 0x11A058: rt::start::hfd264fa826df360859x (in ./Programs/Rust/learn)
==3771== by 0x119E35: rt::lang_start::hde2a214462357c7eb9x (in ./Programs/Rust/learn)
==3771==
==3771== LEAK SUMMARY:
==3771==    definitely lost: 0 bytes in 0 blocks

```

memcheck

Memory related errors

- not so common in Rust due to RAI ??
- !! jemalloc's support for Valgrind controversy, use system allocator

Compile with debug symbols and no optimizations:

(runs 10-50 times slower than natively)

valgrind -v --leak-check=full ./prog args..

--track-fds a list of open file descriptors on exit or on request

--trace-children a list of open file descriptors on exit or on request

--xtree-memory=full Produces execution tree detailing which piece of code is responsible for heap memory usage; display via Kcachgrind

valgrind --tool=memcheck --leak-check=full --track-fds=yes --track-origins=yes ./prog

Types of error messages:

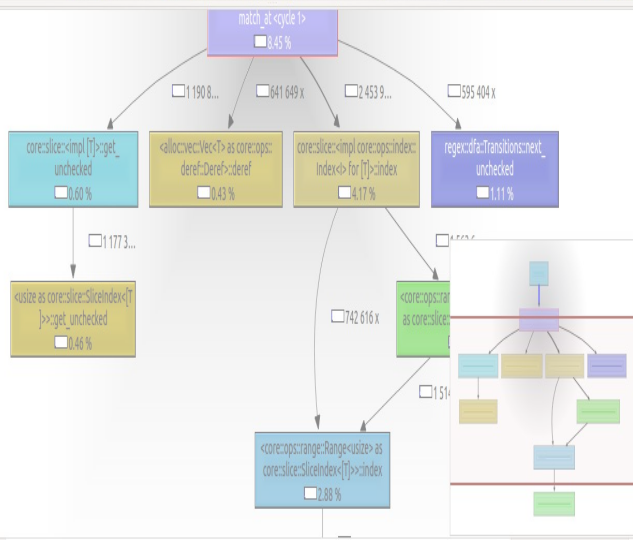
- Definitely lost – memory leaked
- Probably lost – memory leaking unless you do magic with pointers
- Uninitialised values – eg if conditional jump/move depends on uninitialised value

For full list: <http://valgrind.org/docs/manual/mc-manual.html#mc-manual.errormsgs>

Projects memcheck-free of errors (OpenOffice, Firefox,..)

[illegible]

Visualization tool for Valgrind's memcheck, cachegrind, callgrind.



Compile with RUSTFLAGS set with -Z sanitizer flag

(runs ~2x slower than natively)

```
RUSTFLAGS="-Z sanitizer=$SAN" cargo run --target  
x86_64-unknown-linux-gnu prog args
```

Can run with **\$SAN**:

- **address** - detects out of bounds access and use of freed memory
- **leak** - detects memory leaks
- **memory** - reads on uninitialized memory, memory leaks
- **thread**

```

desc: TL cache: 32568 B, 64 B, 8-way, associative
desc: D1 cache: 32568 B, 64 B, 8-way, associative
desc: LL cache: 31457 B, 64 B, 8-way, associative
cmd: ./target/debug/bench
events: Ir 1lMr D1Mr D1Lr D1Ml D1Ml D1Ml D1Ml D1Ml D1Ml D1Ml
fl=/build/glibc-0TSELS/glibc-2.27/csu/.../csu/init-first.c
fn= init
52 8 1 1 0 0 0 0 4 0 0 0 0 0 0 0
55 4 1 1 1 1 0 1 1 0 1 1 1 0 0
62 5 1 1 4 2 0 0 0 0 0 1 1 0 0
67 1 0 0 0 0 0 1 1 1 0 0 0 0 0
68 1 0 0 0 0 0 1 0 0 0 0 0 0 0
69 2 0 0 1 1 0 1 1 1 1 0 0 0 0
81 4 0 0 0 0 0 1 0 0 0 0 0 0 0
84 1 0 0 0 0 0 1 0 0 0 0 0 0 0
89 6 0 0 5 0 0 0 0 0 0 0 0 0 0
fl=/build/glibc-0TSELS/glibc-2.27/csu/.../csu/libc-start.c
fn=(below main)
137 10 2 2 2 0 0 0 7 0 0 0 0 0 0 0
141 4 0 0 1 0 0 1 0 0 1 1 0 0
239 2 0 0 0 0 0 0 0 0 1 0 0 0
240 4 0 0 0 0 0 1 0 0 0 0 0 0
262 4 1 1 2 0 0 0 0 0 1 0 0 0
265 2 0 0 0 0 0 0 0 0 1 0 0 0
266 5 0 0 4 0 0 1 0 0 0 1 1
270 4 0 0 2 0 0 0 0 1 0 0 0 0
285 2 1 1 0 0 0 0 0 1 0 0 0
297 2 0 0 0 0 0 1 0 0 0 0 0
298 2 0 0 0 0 0 0 0 1 0 0 0
303 2 0 0 1 1 0 1 0 0 0 0 0
304 2 0 0 1 0 0 1 0 0 0 0 0
307 2 0 0 0 0 0 1 0 0 0 0 0
310 6 1 1 5 0 0 1 0 0 0 1 1
344 2 1 0 0 0 0 1 0 0 0 0 0
fl=/build/glibc-0TSELS/glibc-2.27/csu/.../sysdeps/generic/dl-
fn= init
43 1 0 0 0 0 0 0 0 0 0 0 0 0 0
44 1 0 0 0 0 0 0 0 0 0 0 0 0
45 13 0 0 4 1 0 0 0 0 4 3 0 0
48 12 0 0 0 0 0 0 0 0 0 0 0 0
62 16 0 0 0 0 0 0 0 0 0 0 0 0
67 1 0 0 0 0 0 0 0 0 0 0 0 0
fl=/build/glibc-0TSELS/glibc-2.27/csu/.../sysdeps/unix/sysv/l
fn= init
36 7 2 2 0 0 0 3 0 0 1 0 0 0
38 4 0 0 0 0 0 1 0 0 0 0 0 0
40 3 0 0 0 0 0 0 0 0 0 0 0 0
41 2 1 1 1 0 0 0 0 0 0 0 0 0
42 1 0 0 0 0 1 1 1 0 0 0 0
44 3 0 0 0 0 0 1 0 0 0 0 0
45 2 0 0 1 0 0 0 0 0 0 0 0 0
46 1 0 0 0 0 1 0 0 0 0 0 0
fl=/build/glibc-0TSELS/glibc-2.27/ctype/ctype-info.c
fn= ctype_init
31 7 2 2 5 3 0 1 0 0 0 0 0 0
33 4 0 0 2 1 0 1 0 0 0 0 0 0
35 4 0 0 2 0 1 0 0 0 0 0 0
36 1 0 0 1 0 0 0 0 0 0 0 0 0
fl=/build/glibc-0TSELS/glibc-2.27/dlfcn/dlerror.c
fn=check_free
186 5 1 1 1 0 0 2 0 0 0 0 0 0

```


Compile with debug symbols and optimizations in:

```
--branch-sim=yes  branch prediction statistics
```

Result displays, eg.:

- ## Cachegrind gathers:

- <http://valgrind.org/docs/manual/cq-manual.html>



Cache misses? ... so what?

Price for cache misses on modern machines:

- L1 miss ~10 cycles
- L2 miss ~15 cycles
- LL miss ~200 cycles
- Mispredicted branch ~10-30 cycles

Detailed branch profiling can help understand how program interacts with the machine and how to make it faster.

Further reading on caches:

<https://www.extremetech.com/extreme/188776-how-l1-and-l2-cpu-caches-work-and-why-theyre-an-essential-part-of-modern-chips>

```

Cache: 32768 B, 64 B, 8-way associative
D1 cache: 32768 B, 64 B, 8-way associative
LL cache: 32768 B, 64 B, 12-way associative
Command: ./build/debug/example-callgrind
Data file: ./build/output/example-callgrind
Events recorded: Ir IImr ILmr Dr DImr DImr Dw DImw DImw Bc Bcm Bi Bim
Events shown: Ir IImr ILmr Dr DImr DImr Dw DImw DImw Bc Bcm Bi Bim
Event sort order: Ir IImr ILmr Dr DImr DImr Dw DImw DImw Bc Bcm Bi Bim
Thresholds: 0.1 100 100 100 100 100 100 100 100 100 100 100
Include dirs:
User annotate:
Auto-annotation:

-----
Ir IImr ILmr Dr DImr DImr Dw DImw DImw Bc Bcm Bi Bim
0,092,578 1,700 3,368 4,213 2,000 0 0 0 133,026 9,608 4,624 455 PROGRAM TOTALS

-----
Ir IImr ILmr Dr DImr DImr Dw DImw DImw Bc Bcm Bi Bim file:function
103,457 19 19 37,962 764 370 8,909 20 0 19,400 830 0 0 /build/glibc-OTS/EL5/
glibc-2.27/elf/dl-lookup.c:do_lookup_x
94,752 3 3 26,670 0 0 27,279 0 0 7,371 1,022 0 0 /rustc/
01856ed52c58aa5ba66a015354d1cc69e9779bdf/src/libcore/iter/range.rs:core::iter::range::next
core::iter::traits::iterator::Iterator for core::ops::range::Range<A>::next
62,979 3 3 16,485 0 0 26,208 0 0 4,179 14 0 0 /home/martina/presentation/
prednaska/example-callgrind/src/main.rs:example-callgrind::bar
59,800 58 55 14,644 214 106 0 0 7,207 482 232 177 /build/glibc-OTS/EL5/
glibc-2.27/string/../sysdeps/x86_64/strcmp.S:strcmp
59,300 9 9 14,225 1,072 826 17 1 0 13,294 1,996 2 2 /build/glibc-OTS/EL5/
glibc-2.27/elf/dl-addr.c:dl_addr
54,264 2 2 22,344 0 0 19,152 0 0 3,192 11 0 0 /rustc/
01856ed52c58aa5ba66a015354d1cc69e9779bdf/src/libcore/iter/range.rs:i32 as
core::iter::range::Step::add_size
52,844 40 40 13,052 13 7 6,470 6 0 10,412 195 288 4 /build/glibc-OTS/EL5/
glibc-2.27/stdio-common/vfscanf.c:IO_vfscanf
47,880 1 1 15,960 0 0 15,960 0 0 0 0 0 0 /rustc/
01856ed52c58aa5ba66a015354d1cc69e9779bdf/src/libcore/intrinsics.rs:core::intrinsics::copy_nonoverlapping
46,409 11 11 9,421 156 121 5,061 11 0 5,757 364 0 0 /build/glibc-OTS/EL5/
glibc-2.27/elf/dl-lookup.c:dl_lookup_symbol_x
46,284 3 3 11,172 0 0 19,152 0 0 1,596 0 0 0 /rustc/
01856ed52c58aa5ba66a015354d1cc69e9779bdf/src/libcore/ptr.rs:core::ptr::swap_nonoverlapping_one
45,672 5 5 9,025 1 1 6,570 2 2 12,906 36 0 0 /build/glibc-OTS/EL5/
glibc-2.27/string/../sysdeps/x86_64/multiarch/memmove-vec-unaligned-erms.S:_memcopy_avx_unaligned_erms
31,920 1 1 11,172 0 0 9,576 0 0 1,596 15 0 0 /rustc/
01856ed52c58aa5ba66a015354d1cc69e9779bdf/src/libcore/num/mod.rs:core::num::ptr_try_from_impls::<impl
core::convert::TryFrom<usize> for u32>::try_from
31,012 23 23 7,369 879 826 3,458 387 361 6,625 68 271 57 /build/glibc-OTS/EL5/
glibc-2.27/elf/../sysdeps/x86_64/dl-machine.h:dl_relocate_object
30,996 1 1 12,915 0 0 5,166 0 0 0 0 0 0 /rustc/
01856ed52c58aa5ba66a015354d1cc69e9779bdf/src/libcore/cmp.rs:core::cmp::impls::<impl
core::cmp::PartialOrd for i32>::lt
23,614 10 10 2,596 5 4 960 1 0 6,488 220 0 0 /build/glibc-OTS/EL5/
glibc-2.27/stdlib/../stdlib/strtol_l.c:_strtol_l_internal
21,713 2 2 5,623 26 0 3,223 3 1 3,220 62 0 0 /build/glibc-OTS/EL5/
glibc-2.27/elf/dl-misc.c:dl_name_match_p
21,602 7 7 4,171 32 32 8 2 2 6,787 1,084 0 0 /build/glibc-OTS/EL5/
glibc-2.27/elf/dl-tunables.c:_GI_tunables_init
20,748 2 2 12,768 0 0 7,980 0 0 0 0 0 0 /rustc/
01856ed52c58aa5ba66a015354d1cc69e9779bdf/src/libcore/intrinsics

```

cg_annotate:

Display results from Cachegrind files.

- global level, with configurable threshold
- line by line counts for specific file

cg_merge:

Merge multiple profile files; might be useful to aggregate cost over multiple runs of the same program.

cg_diff:

Find differences between two profiles; measure how change to a program affected its performance.

callgrind

Records the call history among functions in a program's run as a call-graph.

- Callgrind shows inclusive cost, where the cost of each function includes the cost of all functions it called directly or indirectly
- Built on top of cachegrind
- Platform dependant (does not work on eg. ARM)

Compile with debug symbols and optimizations on:

valgrind --tool=callgrind [callgrind options] ./prog args

Textual output via callgrind_annotate, callgrind_control.

!! Only measures CPU time, so sleeping times are not included. This makes it unsuitable for programs that wait a significant amount of time for network or disk operations to complete.

```
==4119== Events      : Ir Dr Dw IImr DImr DImw IImr DImr DImw Bc Bcm Bi Bim
==4119== Collected : 6087752338 1877217856 1283645182 49600045 1831393 510972 236
80771 244544744 2779086 223545681 4274964
==4119==
==4119== I   refs:      6,087,752,338
==4119== I1 misses:     49,600,045
==4119== L1i misses:     23,601
==4119== I1 miss rate:      0.81%
==4119== L1i miss rate:    0.00%
==4119==
==4119== D   refs:      3,160,863,038 (1,877,217,856 rd + 1,283,645,182 wr)
==4119== D1 misses:     2,342,365 ( 1,831,393 rd +   510,972 wr)
==4119== L1d misses:     98,191 (   17,420 rd +    80,771 wr)
==4119== D1 miss rate:      0.1% (    0.1% +    0.0% )
==4119== L1d miss rate:    0.0% (    0.0% +    0.0% )
==4119==
==4119== LL refs:       51,942,410 ( 51,431,438 rd +    510,972 wr)
==4119== LL misses:      121,792 (  41,021 rd +    80,771 wr)
==4119== LL miss rate:      0.0% (    0.0% +    0.0% )
==4119==
==4119== Branches:    468,090,425 ( 244,544,744 cond + 223,545,681 ind)
==4119== Mispredicts:   7,054,050 (  2,779,086 cond +  4,274,964 ind)
==4119== Mispred rate:    1.5% (    1.1% +    1.9% )
```

```
cmd: ~/(none)
cmd: ~/riprep/target/debug/rg {?:^[a-zA-Z0-9]{3}_?.*\h ~
```

Massif

```
time_unit: i
#-----
snapshot=0
#-----
time=0
mem_heap_B=0
mem_heap_extra_B=0
mem_stacks_B=0
heap_tree=empty
#-----
snapshot=1
#-----
time=69761158
mem_heap_B=913437
mem_heap_extra_B=41867
mem_stacks_B=0
heap_tree=empty
#-----
snapshot=2
#-----
time=118010624
mem_heap_B=931649
mem_heap_extra_B=42847
mem_stacks_B=0
heap_tree=empty
#-----
snapshot=3
#-----
time=190478033
mem_heap_B=1343660
mem_heap_extra_B=61580
mem_stacks_B=0
heap_tree=empty
#-----
snapshot=4
#-----
time=247064033
mem_heap_B=1366236
mem_heap_extra_B=65028
mem_stacks_B=0
heap_tree=detailed
```

```
n12: 1366236 (heap allocation functions) malloc/new/new[], --alloc-fns, etc.
n4: 510948 0x4041C4: ZN59 $LT$alloc..alloc..Global$u20$as$u20$core..alloc..Alloc$GT
$alloc17hb224319826b5958aE.llvm.13452769670659097226 (alloc.rs:151)
n1: 255640 0x4067D0: alloc::raw_vec::RawVec<T,A>::reserve_internal (raw_vec.rs:668)
n1: 255640 0x409530: alloc::raw_vec::RawVec<T,A>::reserve (raw_vec.rs:491)
n1: 255640 0x43A13B: <alloc::vec::Vec<T> as alloc::vec::SpecExtend<T,I>>::spec_extend (vec.rs:1848)
n1: 255640 0x43A308: <alloc::vec::Vec<T> as alloc::vec::SpecExtend<T,I>>::from_iter (vec.rs:1835)
n1: 255640 0x43EFES: <alloc::vec::Vec<T> as
core::iter::traits::collect::FromIterator<T>>::from_iter (vec.rs:1721)
n1: 255640 0x43D03E: core::iter::traits::iterator::Iterator::collect (iterator.rs:1465)
n2: 255640 0x440CED: regex::compile::Compiler::compile_finish (compile.rs:200)
n1: 236560 0x4403A0: ZN5regex7compile8Compiler11compile_one17h57219262258bca28E.llvm.
12913213845200835544 (compile.rs:157)
n3: 236560 0x43FF60: regex::compile::Compiler::compile (compile.rs:129)
n2: 140160 0x4749C1: regex::exec::ExecBuilder::build (exec.rs:313)
n1: 128000 0x46ED7E: regex::re_builder::unicode::RegexBuilder::build (re_builder.rs:79)
n1: 128000 0x46B150: regex::re_unicode::Regex::new (re_unicode.rs:176)
n1: 128000 0x3BE232: core::ops::function::FnOnce::call_once (types.rs:661)
```

Measures heap memory usage (optional stack usage).
Display how much and where the memory was allocated.

Yes: malloc, calloc, realloc, memalign, new, new[]

No: memory allocated with low-level system calls like
mmap, mremap, brk (pages-as-heap=yes, includes also stack)


Compile with debug symbols, optimisations do not matter
valgrind --tool=massif --threshold=0.01 ./prog args

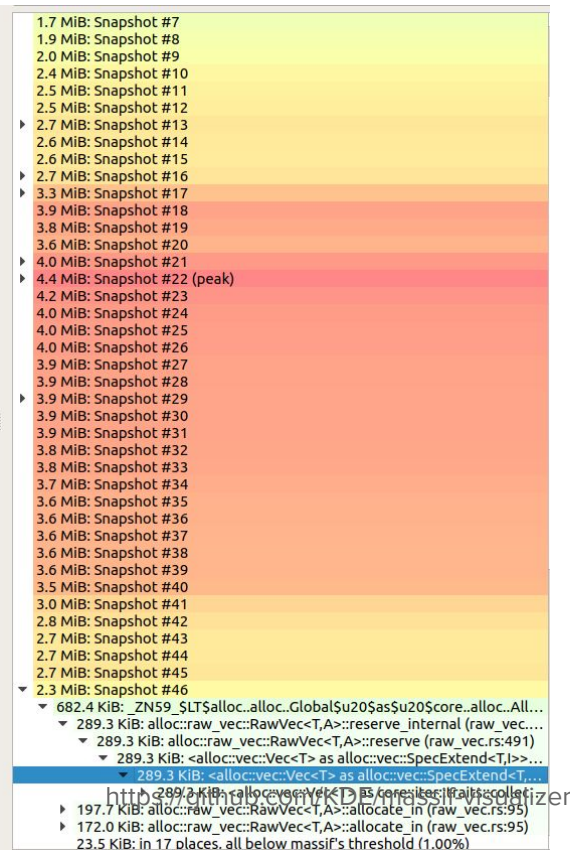

Benefits:

- Reduce the amount of memory program uses → speed up program, helps avoid caching and prevents memory swapping
- Detection of memory leaks

massif_visualizer

Tool that visualizes massif data

- locations that significantly contribute to overall memory consumption
 - memory leaks
 - expensive peaks
- 
- A screenshot of a memory analysis tool's output. It shows a horizontal bar chart where the longest bar is highlighted in green and labeled '1.7 MiB: Snapshot #7'. The chart is part of a larger interface with various tabs and data columns.



```

==3771== HEAP SUMMARY:
==3771==    in use at exit: 1,184 bytes in 4 blocks
==3771== total heap usage: 18 allocs, 14 frees, 2,696 bytes allocated
==3771==
==3771== 32 bytes in 1 blocks are still reachable in loss record 1 of 4
==3771== at 0x128A12: je_malloc (in ./Programs/Rust/learn)
==3771== by 0x11137F: thread local::imp:register_dtor:hb0d4515a39f84cf203b (in ./Programs/Rust/learn)
==3771== by 0x1129E3: thread local::KeySLT9T5GTs::with:h7241079596358413660 (in ./Programs/Rust/learn)
==3771== by 0x112C42: io::stdio::with task stdout:hafb002ec3975fdc5syg (in ./Programs/Rust/learn)
==3771== by 0x115557: io::stdio::println args:h7f0794513d4f2dfcx0g (in ./Programs/Rust/learn)
==3771== by 0x1110B7: main:hb0d1322dca53aee9eaa (in ./Programs/Rust/learn)
==3771== by 0x11A22A: rt::start::closure.32122 (in ./Programs/Rust/learn)
==3771== by 0x124FEB: rust try inner (in ./Programs/Rust/learn)
==3771== by 0x124FD5: rust try (in ./Programs/Rust/learn)
==3771== by 0x1238B2: unwind::try:h72f7fc7ecc46c43Tyc (in ./Programs/Rust/learn)
==3771== by 0x123A88: task::Task::run:h911f3b3bbb0c433efKb (in ./Programs/Rust/learn)
==3771== by 0x11A058: rt::start::hfd264fa826df360859x (in ./Programs/Rust/learn)
==3771==
==3771== 64 bytes in 1 blocks are still reachable in loss record 2 of 4
==3771== at 0x128A12: je_malloc (in ./Programs/Rust/learn)
==3771== by 0x111431: thread local::imp:register_dtor:hb0d4515a39f84cf203b (in ./Programs/Rust/learn)
==3771== by 0x1129E3: thread local::KeySLT9T5GTs::with:h7241079596358413660 (in ./Programs/Rust/learn)
==3771== by 0x112C42: io::stdio::with task stdout:hafb002ec3975fdc5syg (in ./Programs/Rust/learn)
==3771== by 0x115557: io::stdio::println args:h7f0794513d4f2dfcx0g (in ./Programs/Rust/learn)
==3771== by 0x1110B7: main:hb0d1322dca53aee9eaa (in ./Programs/Rust/learn)
==3771== by 0x11A22A: rt::start::closure.32122 (in ./Programs/Rust/learn)
==3771== by 0x124FEB: rust try inner (in ./Programs/Rust/learn)
==3771== by 0x124FD5: rust try (in ./Programs/Rust/learn)
==3771== by 0x1238B2: unwind::try:h72f7fc7ecc46c43Tyc (in ./Programs/Rust/learn)
==3771== by 0x123A88: task::Task::run:h911f3b3bbb0c433efKb (in ./Programs/Rust/learn)
==3771== by 0x11A058: rt::start::hfd264fa826df360859x (in ./Programs/Rust/learn)
==3771==
==3771== 64 bytes in 1 blocks are still reachable in loss record 3 of 4
==3771== at 0x128A12: je_malloc (in ./Programs/Rust/learn)
==3771== by 0x112CD0: io::stdio::with task stdout:hafb002ec3975fdc5syg (in ./Programs/Rust/learn)
==3771== by 0x115557: io::stdio::println args:h7f0794513d4f2dfcx0g (in ./Programs/Rust/learn)
==3771== by 0x1110B7: main:hb0d1322dca53aee9eaa (in ./Programs/Rust/learn)
==3771== by 0x11A22A: rt::start::closure.32122 (in ./Programs/Rust/learn)
==3771== by 0x124FEB: rust try inner (in ./Programs/Rust/learn)
==3771== by 0x124FD5: rust try (in ./Programs/Rust/learn)
==3771== by 0x1238B2: unwind::try:h72f7fc7ecc46c43Tyc (in ./Programs/Rust/learn)
==3771== by 0x123A88: task::Task::run:h911f3b3bbb0c433efKb (in ./Programs/Rust/learn)
==3771== by 0x11A058: rt::start::hfd264fa826df360859x (in ./Programs/Rust/learn)
==3771== by 0x119E35: rt::lang_start:hde2a214462357c7eb9x (in ./Programs/Rust/learn)
==3771== by 0x11117E: main (in ./Programs/Rust/learn)
==3771==
==3771== 1,024 bytes in 1 blocks are still reachable in loss record 4 of 4
==3771== at 0x128A12: je_malloc (in ./Programs/Rust/learn)
==3771== by 0x112222: io::buffered:BufferedWriterLT9W4GTs::with capacity:h8206835914802320726 (in ./Programs/Rust/learn)
==3771== by 0x112063: io::stdio::with task stdout:hafb002ec3975fdc5syg (in ./Programs/Rust/learn)
==3771== by 0x115557: io::stdio::println args:h7f0794513d4f2dfcx0g (in ./Programs/Rust/learn)
==3771== by 0x1110B7: main:hb0d1322dca53aee9eaa (in ./Programs/Rust/learn)
==3771== by 0x11A22A: rt::start::closure.32122 (in ./Programs/Rust/learn)
==3771== by 0x124FEB: rust try inner (in ./Programs/Rust/learn)
==3771== by 0x124FD5: rust try (in ./Programs/Rust/learn)
==3771== by 0x1238B2: unwind::try:h72f7fc7ecc46c43Tyc (in ./Programs/Rust/learn)
==3771== by 0x123A88: task::Task::run:h911f3b3bbb0c433efKb (in ./Programs/Rust/learn)
==3771== by 0x11A058: rt::start::hfd264fa826df360859x (in ./Programs/Rust/learn)
==3771== by 0x119E35: rt::lang_start:hde2a214462357c7eb9x (in ./Programs/Rust/learn)
==3771==
==3771== LEAK SUMMARY:
==3771==    1,024 bytes in 1 blocks are still reachable in loss record 4 of 4

```

Perf (perf events)

Event oriented observability tool which can help you solve advanced performance and troubleshooting functions.

Compile with debug symbols and frame pointers (build with **-fno-omit-frame-pointers**; should be when debug on??).

For full list of all known events: (perf list)

- Event counting (perf stat)
- Profiling (perf record or perf top)
- Static tracing (perf record)
- Dynamic tracing (perf probe)
- Reporting (perf report)
- Cache-2-cache and chacheline false sharing analysis (perf c2c)
- Kernel allocation analysis (perf kmem)
- Lock analysis (perf lock)
- Memory access analysis (perf mem)
- Kernel scheduler analysis (perf sched)

Perf_events Events sources

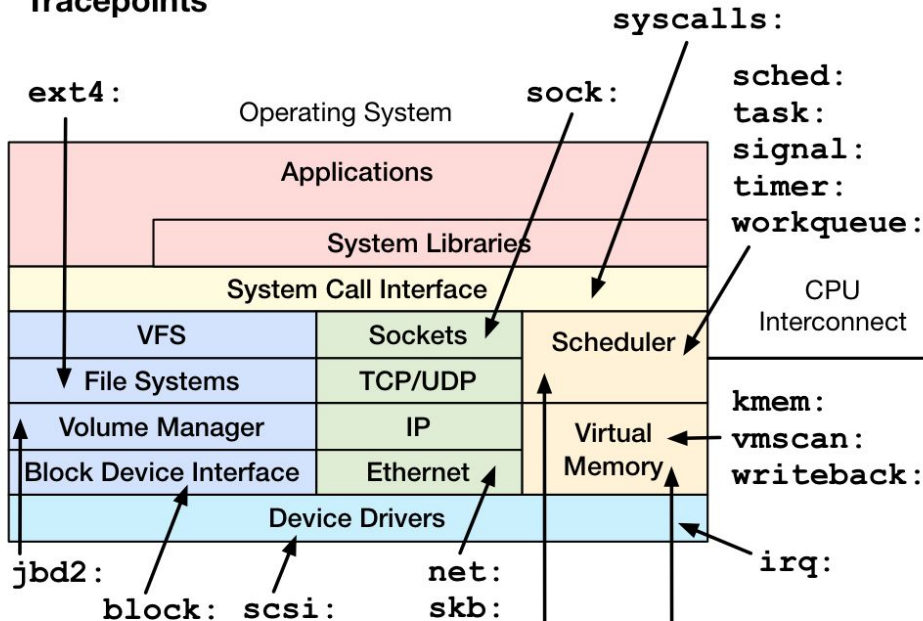
Linux perf_events Event Sources

Dynamic
Tracing

uprobes

kprobes

Tracepoints




Software Events

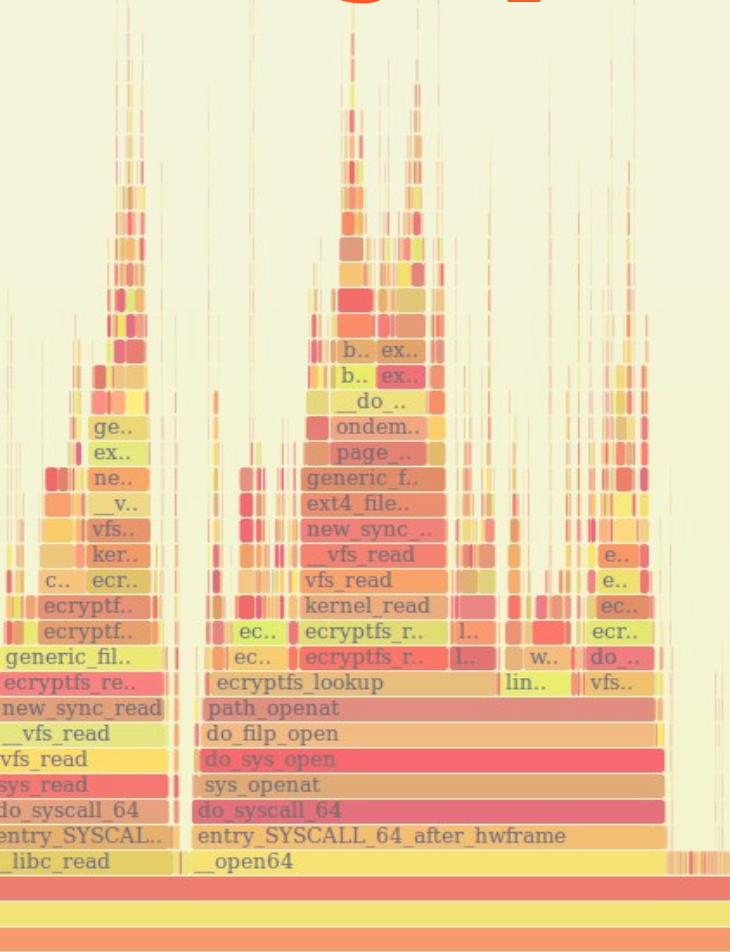
cpu-clock
cs migrations

page-faults
minor-faults
major-faults

PMCs



flamegraph



Helps quantify code-paths and determine places with performance problems

Visualization of perf.data from perf report:

- X axis sample population
- Y axis stack depth
- Width of each function is relative to the number of samples.

1. Capture stack:

```
perf record -F 99 --call-graph=dwarf -p PID  
perf script > out.perf
```

2. Use the stackcollapse programs to fold stack samples into single lines:

```
./stackcollapse-perf.pl out.perf > out.folded
```

3. Use flamegraph.pl to render a SVG.

```
./flamegraph.pl out.kern_folded > kernel.svg
```

FlameScope

- Interactive visualization tool for different timeranges of FlameGraph in the form of heat map
- perf script's output that includes stack traces, including page faults, context switches, and other events.



```
sudo perf record -F _FREQUENCY_ -p PID
sudo perf script --header > flamescope_out_file
```


Google perf tools gperftools

Collection of high-performance multi-threaded malloc() implementation and performance analysis tools.

- uses statistical sampling

Rust project that is using gperf tools - **cpuprofiler**

<https://github.com/AtheMathmo/cpuprofiler>

- Heap profiler (malloc)
- Cpu profiler
- Heap leak-checker

Visualizer tools:

- GraphViz
- pprof

Other tools

CPU bound problems:

toplev CPU performance (counting) tool
<https://github.com/andikleen/pmu-tools/wiki/toplev-manual>

Off-CPU analysis:

eBPF an universal in-kernel virtual machine, that has hooks all over the kernel.

<http://www.brendangregg.com/blog/2019-01-01/learn-ebpf-tracing.html>
<http://www.brendangregg.com/offcpuanalysis.html>

Monitoring framework:

Vector An on-host performance monitoring framework which exposes hand picked high resolution metrics to every engineer's browser.
<https://github.com/Netflix/vector>

Discussion

- Don't try to over optimize the program!
- Kachegrind tip page: When all functions take almost equal time and the program feels fast for you: stop! There is no reason to make the program 1% faster with days of optimization.
- Know what the issue is / narrow down the problem before using any tool!
- Know your program and do some expectations for different parts of the code!
- Do not use random tools!
- Do not tune things at random until problem goes away!
- Do not blame someone else, measure!